

INFS2200 ASSIGNMENT

Student Number: 4536071

Task 1 – Constraints

1.1 Find Existing Constraints

```
SELECT CONSTRAINT_NAME, TABLE_NAME, INDEX_NAME
FROM USER_CONSTRAINTS WHERE TABLE_NAME = 'EMP';
```

CONSTRAINT_NAME	TABLE_NAME	INDEX_NAME
PK_EMPNO	EMP	PK_EMPNO

```
SELECT CONSTRAINT_NAME, TABLE_NAME, INDEX_NAME
FROM USER_CONSTRAINTS WHERE TABLE_NAME = 'DEPT';
```

CONSTRAINT_NAME	TABLE_NAME	INDEX_NAME
PK_DEPTNO	DEPT	PK_DEPTNO

```
SELECT CONSTRAINT_NAME, TABLE_NAME, INDEX_NAME
FROM USER_CONSTRAINTS WHERE TABLE_NAME = 'PURCHASE';
```

CONSTRAINT_NAME	TABLE_NAME	INDEX_NAME
PK_PURCHASENO	PURCHASE	PK_PURCHASENO

```
SELECT CONSTRAINT_NAME, TABLE_NAME, INDEX_NAME
FROM USER_CONSTRAINTS WHERE TABLE_NAME = 'CLIENT';
```

CONSTRAINT_NAME	TABLE_NAME	INDEX_NAME
PK_CLIENTNO	CLIENT	PK_CLIENTNO

1.2 Create Missing Constraints

```
ALTER TABLE DEPT ADD CONSTRAINT UN_DNAME UNIQUE (DName);

ALTER TABLE PURCHASE MODIFY AMOUNT INT CONSTRAINT CK_AMOUNT NOT NULL;
ALTER TABLE EMP MODIFY ENAME VARCHAR(30) CONSTRAINT CK_ENAME NOT NULL;
ALTER TABLE DEPT MODIFY DNAME VARCHAR(30) CONSTRAINT CK_DNAME NOT NULL;
ALTER TABLE CLIENT MODIFY CNAME VARCHAR(30) CONSTRAINT CK_CNAME NOT NULL;
ALTER TABLE PURCHASE MODIFY RECEIPTNO INT CONSTRAINT CK_RECEIPTNO NOT NULL;

ALTER TABLE PURCHASE ADD CONSTRAINT CK_SERVICETYPE
CHECK (ServiceType IN ('Training', 'Data Recovery', 'Consultation',
'Software Installation', 'Software Repair'));

ALTER TABLE PURCHASE ADD CONSTRAINT CK_PAYMENTTYPE
CHECK (PaymentType IN ('Debit', 'Cash', 'Credit'));

ALTER TABLE PURCHASE ADD CONSTRAINT CK_GST
CHECK (GST IN ('Yes', 'No'));

ALTER TABLE EMP ADD CONSTRAINT FK_DEPTNO FOREIGN KEY (DeptNo)
REFERENCES DEPT (DeptNo);

ALTER TABLE PURCHASE ADD CONSTRAINT FK_EMPNO FOREIGN KEY (ServedBy)
REFERENCES EMP (EmpNo);

ALTER TABLE PURCHASE ADD CONSTRAINT FK_CLIENTNO FOREIGN KEY (ClientNo)
REFERENCES CLIENT (ClientNo);
```

CONSTRAINT_NAME	TABLE_NAME	INDEX_NAME
FK_DEPTNO	EMP	
CK_ENAME	EMP	
PK_EMPNO	EMP	PK_EMPNO

CONSTRAINT_NAME	TABLE_NAME	INDEX_NAME
CK_DNAME	DEPT	
UN_DNAME	DEPT	UN_DNAME
PK_DEPTNO	DEPT	PK_DEPTNO

CONSTRAINT_NAME	TABLE_NAME	INDEX_NAME
FK_EMPNO	PURCHASE	
FK_CLIENTNO	PURCHASE	
CK_AMOUNT	PURCHASE	
CK_RECEIPTNO	PURCHASE	
CK_SERVICETYPE	PURCHASE	
CK_PAYMENTTYPE	PURCHASE	
CK_GST	PURCHASE	
PK_PURCHASENO	PURCHASE	PK_PURCHASENO

CONSTRAINT_NAME	TABLE_NAME	INDEX_NAME
CK_CNAME	CLIENT	
PK_CLIENTNO	CLIENT	PK_CLIENTNO

Task 2 – Triggers

2.1 Top Client

```
CREATE TABLE TOPCOMPANYCLIENT AS
SELECT C.ClientNo AS ClientNumber, C.CName AS ClientName,
SUM(P.Amount) AS TotalAmount
FROM CLIENT C, PURCHASE P
WHERE C.ClientNo = P.ClientNo
GROUP BY C.ClientNo, C.CName
ORDER BY SUM(P.Amount) DESC
FETCH FIRST 1 ROWS ONLY;

SELECT * FROM TOPCOMPANYCLIENT;
```

CLIENTNUMBER	CLIENTNAME	TOTALAMOUNT
24535	Sally Moon	20100

2.2 Top Discount Trigger

```
CREATE OR REPLACE TRIGGER TOP_DISCOUNT
BEFORE INSERT ON PURCHASE
FOR EACH ROW
DECLARE
    ClientID INT;
BEGIN
    SELECT T.ClientNumber INTO ClientID
    FROM TOPCOMPANYCLIENT T;
    IF :NEW.CLIENTNO = ClientID THEN
        :NEW.AMOUNT := (0.85 * :NEW.AMOUNT);
    END IF;
END;
/
```

2.2 Trigger Check

```
SELECT *
FROM PURCHASE P
WHERE P.CLIENTNO = 24535
AND P.PAYMENTTYPE = 'Cash'
AND P.GST = 'No'
AND P.ServedBy = 1045;
```

PURCHASENO	RECEIPTNO	SERVICETYPE	PAYMENTTYP	GST	AMOUNT	SERVEDBY	CLIENTNO
346	546443	Software Repair	Cash	No	915	1045	24535

```
INSERT INTO PURCHASE VALUES (10000, 101101, 'Consultation', 'Cash', 'No', 1000, 1045, 24535);
```

PURCHASENO	RECEIPTNO	SERVICETYPE	PAYMENTTYP	GST	AMOUNT	SERVEDBY	CLIENTNO
346	546443	Software Repair	Cash	No	915	1045	24535
10000	101101	Consultation	Cash	No	850	1045	24535

Above, a new purchase made by our top client was inserted into the purchase table with the amount 1000. The trigger works and we can see this by the updated screenshot of the table above with the amount at 850.

2.3 Sales - Sunshine Department

```
CREATE OR REPLACE TRIGGER SUNSHINE_DEPT
BEFORE INSERT ON PURCHASE
FOR EACH ROW
DECLARE
    DepartName VARCHAR(30);
BEGIN
    SELECT DISTINCT D.DName INTO DepartName
    FROM EMP E, DEPT D, PURCHASE P
    WHERE E.DEPTNO = D.DEPTNO
    AND E.EMPNO = P.SERVEDBY
    AND :NEW.SERVEDBY = P.SERVEDBY;

    IF DepartName = 'SALES - Sunshine'
    AND :NEW.PAYMENTTYPE != 'Cash' THEN
        :NEW.PAYMENTTYPE := 'Cash';
    END IF;

    IF DepartName = 'SALES - Sunshine'
    AND :NEW.SERVICETYPE = 'Data Recovery' THEN
        :NEW.AMOUNT := (:NEW.AMOUNT * 0.70);
    END IF;
END;
/
```

2.3 Trigger Check

Check 1: Update Columns to Cash:

```
INSERT INTO PURCHASE VALUES (90000, 522555, 'Consultation', 'Credit', 'Yes', 1000, 7777,
24542);

SELECT *
FROM PURCHASE P
WHERE P.PurchaseNo = 90000;
```

PURCHASENO	RECEIPTNO	SERVICETYPE	PAYMENTTYP	GST	AMOUNT	SERVEDBY	CLIENTNO
90000	522555	Consultation	Cash	Yes	1000	7777	24542

Check 2: Update to Cash and Amount for Data Recovery:

```
INSERT INTO PURCHASE VALUES (80000, 522555, 'Data Recovery', 'Debit', 'No', 1000, 7777, 24542);

SELECT *
FROM PURCHASE P
WHERE P.PurchaseNo = 80000;
```

PURCHASENO	RECEIPTNO	SERVICETYPE	PAYMENTTYP	GST	AMOUNT	SERVEDBY	CLIENTNO
80000	522555	Data Recovery	Cash	No	700	7777	24542

Task 3 – Views

3.1 Virtual View

```
CREATE OR REPLACE VIEW V_DEPT_AMOUNT AS
SELECT D.DEPTNO, D.DNAME, AVG(P.AMOUNT) AS DAVGAMT, MAX(P.AMOUNT) AS DMAXAMT,
MIN(P.AMOUNT) AS DMINAMT, SUM(P.AMOUNT) AS DTOTAL
FROM DEPT D, PURCHASE P, EMP E
WHERE E.DEPTNO = D.DEPTNO
AND E.EMPNO = P.SERVEDBY
GROUP BY D.DEPTNO, D.DNAME;
```

DEPTNO	DNAME	DAVGAMT	DMAXAMT	DMINAMT	DTOTAL
30	SALES - Sunflower	528.336968	1000	50	968970
40	SALES - Hercules	535.761089	1000	50	1062950
50	SALES - Neptune	517.578053	1000	50	1674365
20	SALES - Sunshine	522.126719	1000	50	1063050
10	SALES - Glorious	522.730769	1000	50	475685

Elapsed: 00:00:00.04

3.2 Materialised View

```
CREATE MATERIALIZED VIEW MV_DEPT_AMOUNT
BUILD IMMEDIATE
AS
SELECT D.DEPTNO, D.DNAME, AVG(P.AMOUNT) AS DAVGAMT, MAX(P.AMOUNT) AS DMAXAMT,
MIN(P.AMOUNT) AS DMINAMT, SUM(P.AMOUNT) AS DTOTAL
FROM DEPT D, PURCHASE P, EMP E
WHERE E.DEPTNO = D.DEPTNO
AND E.EMPNO = P.SERVEDBY
GROUP BY D.DEPTNO, D.DNAME;
```

DEPTNO	DNAME	DAVGAMT	DMAXAMT	DMINAMT	DTOTAL
30	SALES - Sunflower	528.336968	1000	50	968970
40	SALES - Hercules	535.761089	1000	50	1062950
50	SALES - Neptune	517.578053	1000	50	1674365
20	SALES - Sunshine	522.126719	1000	50	1063050
10	SALES - Glorious	522.730769	1000	50	475685

Elapsed: 00:00:00.00

3.3 Execution Time

The materialised view did speed up query processing, as can be viewed below:

```
EXPLAIN PLAN FOR SELECT * FROM V_DEPT_AMOUNT;
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);
```

PLAN_TABLE_OUTPUT							
Plan hash value: 58894710							
Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	
0	SELECT STATEMENT		9218	738K	25 (12)	00:00:01	
1	HASH GROUP BY		9218	738K	25 (12)	00:00:01	
2	NESTED LOOPS		9218	738K	23 (5)	00:00:01	
3	NESTED LOOPS		9218	738K	23 (5)	00:00:01	
4	NESTED LOOPS		9218	468K	23 (5)	00:00:01	
5	TABLE ACCESS FULL	PURCHASE	9218	234K	22 (0)	00:00:01	
PLAN_TABLE_OUTPUT							
6	TABLE ACCESS BY INDEX ROWID	EMP	1	26	0 (0)	00:00:01	
* 7	INDEX UNIQUE SCAN	PK_EMPNO	1		0 (0)	00:00:01	
* 8	INDEX UNIQUE SCAN	PK_DEPTNO	1		0 (0)	00:00:01	
9	TABLE ACCESS BY INDEX ROWID	DEPT	1	30	0 (0)	00:00:01	

```
EXPLAIN PLAN FOR SELECT * FROM MV_DEPT_AMOUNT;
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);
```

PLAN_TABLE_OUTPUT							
Plan hash value: 3751359186							
Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	
0	SELECT STATEMENT		5	270	3 (0)	00:00:01	
1	MAT_VIEW ACCESS FULL	MV_DEPT_AMOUNT	5	270	3 (0)	00:00:01	
8 rows selected.							

3.3 Summary:

Clearly, the materialized view sped up the query processing by 00:00:00.04 according to the screenshots of elapsed time above.

Also, the cost is significantly less at 3 for the materialized view whereas the virtual view cost ranges from 25 to 22 depending on the operation. Also, number of bytes was lower at 270 for each materialized view operation compared to 738K for the select statements and nested loop operations for the virtual view.

The number of operations for the materialized view only has 2: one for the select statement and one to process the view. The virtual view has 1 select statement operations and then 9 number operations for the nested loops and other SQL logic. We can also see that the number of rows that were accessed is 9218 for just the select statement for a virtual view in comparison to 5 for the select statement of the materialized view.

3.4 Employee Views

3.4.1 Virtual View

```
CREATE OR REPLACE VIEW V_DEPT_EMP_AMOUNT AS
SELECT E.EMPNO, D.DEPTNO, COUNT(P.PurchaseNo) AS ETOTALPNUM,
AVG(P.AMOUNT) AS EAVGPURCH, MAX(P.AMOUNT) AS EMAXPURCH,
SUM(P.Amount) AS ETOTALSERVED
FROM DEPT D, PURCHASE P, EMP E
WHERE E.DEPTNO = D.DEPTNO
AND E.EMPNO = P.SERVEDBY
GROUP BY E.EMPNO, D.DEPTNO
ORDER BY D.DEPTNO ASC, SUM(P.Amount) DESC;
```

3.4.2 Materialized View

```
CREATE MATERIALIZED VIEW MV_DEPT_EMP_AMOUNT
BUILD IMMEDIATE
AS
SELECT E.EMPNO, D.DEPTNO, COUNT(P.PurchaseNo) AS ETOTALPNUM,
AVG(P.AMOUNT) AS EAVGPURCH, MAX(P.AMOUNT) AS EMAXPURCH,
SUM(P.Amount) AS ETOTALSERVED
FROM DEPT D, PURCHASE P, EMP E
WHERE E.DEPTNO = D.DEPTNO
AND E.EMPNO = P.SERVEDBY
GROUP BY E.EMPNO, D.DEPTNO
ORDER BY SUM(P.Amount) DESC, D.DEPTNO ASC;
```

Virtual View:

```
SELECT * FROM V_DEPT_EMP_AMOUNT;
```

Materialized View:

```
SELECT * FROM MV_DEPT_EMP_AMOUNT;
```


View output:

```
SQL> SELECT * FROM V_DEPT_EMP_AMOUNT;
```

EMPNO	DEPTNO	ETOTALPNUM	EAVGPURCH	EMAXPURCH	ETOTALSERVED
1065	10	166	543.614458	990	90240
1007	10	167	502.664671	990	83945
1015	10	150	544.233333	995	81635
1031	10	144	524.305556	1000	75500
1009	10	138	546.195652	1000	75375
1055	10	144	472.638889	985	68060
1071	10	1	930	930	930
1022	20	167	533.952006	995	89170
1049	20	155	550.548387	1000	85335
1039	20	158	529.683544	995	83690
1017	20	146	543.390411	1000	79335
1028	20	145	544.827586	995	79080
1036	20	142	545.387324	995	77445
1002	20	144	525.243056	1000	75635
1006	20	156	480.737179	975	74995
1037	20	141	529.893617	985	74715
1010	20	146	503.184932	1000	73465
1013	20	139	502.769784	1000	69885
1058	20	142	491.795775	995	69835
1040	20	137	508.029197	995	69600
1060	20	117	519.829006	975	60820
1072	20	1	125	125	125
1020	30	167	543.323353	995	90735
1067	30	168	529.22619	1000	88910
1048	30	150	571.433333	990	85715
1061	30	141	560.496454	1000	79030
1044	30	147	534.931973	995	78635
1001	30	152	504.638158	995	76705
1069	30	147	517.823129	1000	76120
1027	30	126	545.793651	1000	68770
1053	30	141	486.737589	995	68030
1052	30	131	521.030534	990	60255
1064	30	128	527.03125	995	67460
1066	30	123	505.731707	995	62205
1062	30	112	511.651786	990	57305
1073	30	1	495	495	495
1019	40	168	533.928571	1000	89700
1021	40	155	575.064516	1000	89135
1059	40	152	571.546053	995	86075
1012	40	150	545.466667	1000	81820
1003	40	150	531.766667	985	79765
1004	40	135	584.481481	995	78905
1005	40	146	506.335616	990	73925
1042	40	148	480.641892	995	71135

3.5 Execution Time:

Virtual View:

```
EXPLAIN PLAN FOR SELECT * FROM V_DEPT_EMP_AMOUNT;  
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);
```

```
75 rows selected.  
Elapsed: 00:00:00.05
```

```
SQL> SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);
```

PLAN_TABLE_OUTPUT						
Plan hash value: 740344440						
Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		9218	702K	25 (12)	00:00:01
1	VIEW	V_DEPT_EMP_AMOUNT	9218	702K	25 (12)	00:00:01
2	SORT ORDER BY		9218	819K	25 (12)	00:00:01
3	NESTED LOOPS		9218	819K	24 (9)	00:00:01
4	NESTED LOOPS		9218	819K	24 (9)	00:00:01
5	VIEW	VW_GBC_6	9218	585K	23 (5)	00:00:01
PLAN_TABLE_OUTPUT						
6	HASH GROUP BY		9218	234K	23 (5)	00:00:01
7	TABLE ACCESS FULL	PURCHASE	9218	234K	22 (0)	00:00:01
* 8	INDEX UNIQUE SCAN	PK_EMPNO	1	0	0 (0)	00:00:01
* 9	TABLE ACCESS BY INDEX ROWID	EMP	1	26	0 (0)	00:00:01

Materialized View:

```
EXPLAIN PLAN FOR SELECT * MV_DEPT_EMP_AMOUNT;  
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);
```

```
75 rows selected.  
  
Elapsed: 00:00:00.02
```

```
PLAN_TABLE_OUTPUT  
-----  
Plan hash value: 1444910117  
  
-----  
| Id | Operation                | Name                | Rows | Bytes | Cost (%CPU)| Time     |  
-----  
|  0 | SELECT STATEMENT          |                     |    75 | 3000 |      3  (0)| 00:00:01 |  
|  1 |   MAT_VIEW ACCESS FULL    | MV_DEPT_EMP_AMOUNT |    75 | 3000 |      3  (0)| 00:00:01 |  
-----
```

3.5 Summary:

Clearly, the materialized view sped up the query processing by 00:00:00.03 according to the screenshots of elapsed time above. Also, the cost is significantly less at 3 for the materialized view whereas the virtual view cost ranges from 25 to 22 depending on the operation. We can also see that the number of bytes was lower at 300 for each materialized operation compared to 702K for the select statement for the virtual view. The number of operations for the materialized view only has 2: one for the select statement and one to process the view. The virtual view has 1 select statement operations and then 9 number operations for the nested loops and other SQL logic.

Task 4 – Indexes

4.1 Number of purchases

```
SELECT COUNT(*) AS RECEIPT_BOOKS  
FROM (SELECT TRUNC(SUBSTR(P2.ReceiptNo, 1, 3)) AS Books,  
        (SELECT COUNT(1)  
         FROM PURCHASE P  
         WHERE TRUNC(SUBSTR(P.ReceiptNo, 1, 3)) =  
               TRUNC(SUBSTR(P2.ReceiptNo, 1, 3))) AS NumOfBooks  
 FROM PURCHASE P2) Specified  
WHERE Specified.NumOfBooks >= 10;
```

```
RECEIPT_BOOKS  
-----  
7896  
  
Elapsed: 00:00:41.85
```

4.2 Function-based Index

Setting up a plan for query:

```
EXPLAIN PLAN FOR SELECT COUNT(*) AS RECEIPT_BOOKS
FROM (SELECT TRUNC(SUBSTR(P2.ReceiptNo, 1, 3)) AS Books,
      (SELECT COUNT(1)
       FROM PURCHASE P
       WHERE TRUNC(SUBSTR(P.ReceiptNo, 1, 3)) =
            TRUNC(SUBSTR(P2.ReceiptNo, 1, 3))) AS NumOfBooks
FROM PURCHASE P2) Specified
WHERE Specified.NumOfBooks >= 10;
```

Execution Plan without an Index:

```
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);
```

PLAN_TABLE_OUTPUT

Plan hash value: 2870517031

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	13	6498 (3)	00:00:01
1	SORT AGGREGATE		1	13		
* 2	FILTER					
3	TABLE ACCESS FULL	PURCHASE	9218	117K	22 (0)	00:00:01
4	SORT AGGREGATE		1	13		
* 5	TABLE ACCESS FULL	PURCHASE	92	1196	22 (0)	00:00:01

PLAN_TABLE_OUTPUT

Now by creating an index on this table:

```
CREATE INDEX BOOK_INDEX ON PURCHASE (TRUNC(SUBSTR(ReceiptNo, 1, 3)));
```

```
RECEIPT_BOOKS
-----
          7896

Elapsed: 00:00:00.04
SQL>
```

Execution Plan with Index in Place:

PLAN_TABLE_OUTPUT

Plan hash value: 589993641

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	13	310 (0)	00:00:01
1	SORT AGGREGATE		1	13		
* 2	FILTER					
3	TABLE ACCESS FULL	PURCHASE	9218	117K	22 (0)	00:00:01
4	SORT AGGREGATE		1	13		
* 5	INDEX RANGE SCAN	BOOK_INDEX	92	1196	1 (0)	00:00:01

PLAN_TABLE_OUTPUT

With the index in place, the query speed up incredibly. As we can see it took the query 41.85 seconds without the index in place whereas when an index was created on the table, it only took 00.04 seconds to process. This is a significant difference in execution time. The cost is also less at $310 + 22 + 1 = 333$ total with an index and $6498 + 22 + 22 = 6542$ without an index. Therefore a cost difference of 6209. Therefore using an index is a great decision for this query.

```
SELECT SUM(TOTAL)
FROM (SELECT P.AMOUNT AS TOTAL
      FROM PURCHASE P, DEPT D, EMP E
      WHERE E.EMPNO = P.SERVEDBY
      AND D.DEPTNO = E.DEPTNO
      AND D.DEPTNO = 50
      AND P.PURCHASENO NOT IN (SELECT P2.PURCHASENO
                                FROM PURCHASE P2
                                WHERE INSTR(P2.ServiceType, 'Software') > 0));
```

```
EXPLAIN PLAN FOR SELECT SUM(TOTAL)
FROM (SELECT P.AMOUNT AS TOTAL
      FROM PURCHASE P, DEPT D, EMP E
      WHERE E.EMPNO = P.SERVEDBY
      AND D.DEPTNO = E.DEPTNO
      AND D.DEPTNO = 50
      AND P.PURCHASENO NOT IN (SELECT P2.PURCHASENO
                              FROM PURCHASE P2
                              WHERE INSTR(P2.ServiceType, 'Software') > 0));

SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);
```

```

PLAN_TABLE_OUTPUT
-----
Plan hash value: 1715138504

-----
| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
-----
| 0 | SELECT STATEMENT | | 1 | 66 | 23 (5) | 00:00:01 | |
| 1 | SORT AGGREGATE | | 1 | 66 | | | |
| 2 | NESTED LOOPS | | 1588 | 102K | 23 (5) | 00:00:01 |
| 3 | NESTED LOOPS | | 4962 | 102K | 23 (5) | 00:00:01 |
| * 4 | TABLE ACCESS FULL | PURCHASE | 4962 | 193K | 22 (0) | 00:00:01 |
| * 5 | INDEX UNIQUE SCAN | PK_EMPNO | 1 | | 0 (0) | 00:00:01 |
-----

PLAN_TABLE_OUTPUT
-----
| * 6 | TABLE ACCESS BY INDEX ROWID | EMP | 1 | 26 | 0 (0) | 00:00:01 |
-----

```

4.4 Service Index

```
CREATE INDEX SERVICE_INDEX ON PURCHASE(INSTR(ServiceType, 'Software'));
```

```
SUM(TOTAL)
-----
905355
Elapsed: 00:00:00.02
```

PLAN_TABLE_OUTPUT							
Plan hash value: 3392947326							
Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	
0	SELECT STATEMENT		1	91	23 (5)	00:00:01	
1	SORT AGGREGATE		1	91			
2	NESTED LOOPS ANTI		2950	262K	23 (5)	00:00:01	
3	NESTED LOOPS		2950	187K	23 (5)	00:00:01	
4	TABLE ACCESS FULL	PURCHASE	9218	351K	22 (0)	00:00:01	
* 5	TABLE ACCESS BY INDEX ROWID	EMP	1	26	0 (0)	00:00:01	
PLAN_TABLE_OUTPUT							
* 6	INDEX UNIQUE SCAN	PK_EMPNO	1		0 (0)	00:00:01	
* 7	TABLE ACCESS BY INDEX ROWID	PURCHASE	1	26	0 (0)	00:00:01	
* 8	INDEX UNIQUE SCAN	PK_PURCHASENO	1		0 (0)	00:00:01	

4.4 Summary:

This service index only just speeds up the query. The cost is the same for both with and without the index implemented. Therefore, implementing an index does not really help for this particular query.

4.5 Number of purchases with same PType, STYPE and GST

```
SELECT SUM(NUMPURCH) AS Number_of_Purchases
FROM (SELECT COUNT(P.PurchaseNo) AS NUMPURCH
      FROM PURCHASE P
      GROUP BY P.ServiceType, P.PaymentType, P.GST
      HAVING COUNT(P.PurchaseNo) >= 1000);
```

```
NUMBER_OF_PURCHASES
-----
2202
```

4.6 Index on 4.5 Query

A non-clustered index would be the best choice for this query. Non-clustered indexes do not sort the physical data inside the table. A non-clustered index locates the table content in one place and the index is located in another. This is what allows more than one index to be applied to a table.

In particular, bitmap Indexes would be the best option because they are better on finite fields. They work best for columns that have a low cardinality. As GST contains Boolean data: "Yes" or "No" options and the other columns, ServiceType and PaymentType, also only have a few options, these tables have a low cardinality.

```
CREATE BITMAP INDEX PURCHASENO_INDEX ON PURCHASE(ServiceType);
CREATE BITMAP INDEX PURCHASENO_INDEX ON PURCHASE(PaymentType);
CREATE BITMAP INDEX PURCHASENO_INDEX ON PURCHASE(GST);
```

Task 5 – Execution Plan

5.1 Purchase number Oracle Optimisation:

```
EXPLAIN PLAN FOR SELECT *  
FROM PURCHASE P  
WHERE P.PurchaseNo = 1234;  
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);
```

```
PLAN_TABLE_OUTPUT  
-----  
Plan hash value: 2822030489  
  
-----  
| Id | Operation          | Name          | Rows | Bytes | Cost (%CPU)| Time |  
-----  
|  0 | SELECT STATEMENT   |               |     1 |    89 |     0 (0)| 00:00:01 |  
|  1 | TABLE ACCESS BY INDEX ROWID | PURCHASE      |     1 |    89 |     0 (0)| 00:00:01 |  
|*  2 | INDEX UNIQUE SCAN  | PK_PURCHASENO |     1 |      |     0 (0)| 00:00:01 |  
-----  
  
Predicate Information (identified by operation id):  
-----  
PLAN_TABLE_OUTPUT  
-----  
  
      2 - access("P"."PURCHASENO">=1234)  
  
14 rows selected.
```

In this execution plan, the index on PK_PURCHASENO is used in a unique scan operation to evaluate the WHERE clause. It then returns a rowID from the index. So, then rows are located by index (TABLE ACCESS BY INDEX ROWID) which means that the full table does not have to be accessed. The SELECT only returns rows satisfying the WHERE clause. From the plan, we can see that 89 bytes were accessed for both the SELECT and TABLE ACCESS BY INDEX ROWID operations.

5.2 Dropping Constraint affects Execution Plan

```
Drop Primary Key Constraint:  
ALTER TABLE PURCHASE  
DROP CONSTRAINT PK_PURCHASENO;
```

```
EXPLAIN PLAN FOR SELECT *  
FROM PURCHASE P  
WHERE P.PurchaseNo = 1234;  
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);
```

```
PLAN_TABLE_OUTPUT  
-----  
Plan hash value: 2913724801  
  
-----  
| Id | Operation          | Name          | Rows | Bytes | Cost (%CPU)| Time |  
-----  
|  0 | SELECT STATEMENT   |               |     1 |    89 |    22 (0)| 00:00:01 |  
|*  1 | TABLE ACCESS FULL | PURCHASE      |     1 |    89 |    22 (0)| 00:00:01 |  
-----  
  
Predicate Information (identified by operation id):  
-----  
PLAN_TABLE_OUTPUT  
-----  
  
      1 - filter("P"."PURCHASENO">=1234)  
  
Note  
-----  
      - dynamic statistics used: dynamic sampling (level=2)  
  
17 rows selected.
```

In this execution plan, every row in the PURCHASE table is accessed and the WHERE clause condition is checked and dynamic sampling was used. One difference between this plan and the 5.1 plan is that this query selects 17 rows rather than 14. Also, the Cost for this execution plan is 22 for both operations whereas with the constraint in place, there is no cost for any operations.